

AI_Assgn_2_Q4

November 10, 2021

1 Import the packages

```
[ ]: import torch
import torchvision
from torch import nn, optim
from torchsummary import summary
```

2 Declare variables for the CNN

- **Epoch** is the number of passes of the entire training dataset through the neural network. A pair of forward and backward propagation indicates a single pass.
- **Batch Size** is the number of samples to work through before updating the weights and biases associated with the model.
- **Learning Rate** controls how much to change the model parameters in response to the prediction error each time the model weights are updated.

```
[ ]: batch_size = 32
epoch = 30
learning_rate = 0.01
```

3 Load the training set and validation set using Dataset and DataLoader

```
[ ]: trans = torchvision.transforms.ToTensor()
train_data = torch.utils.data.DataLoader(
    torchvision.datasets.MNIST(
        'mnist_data', train=True, download=True, transform=trans
    ), batch_size=batch_size
)
val_data = torch.utils.data.DataLoader(
    torchvision.datasets.MNIST(
        'mnist_data', train=False, download=True, transform=trans
    ), batch_size=batch_size)
```

Downloading http://yann.lecun.com/exdb/mnist/train-images-idx3-ubyte.gz
Downloading http://yann.lecun.com/exdb/mnist/train-images-idx3-ubyte.gz to
mnist_data/MNIST/raw/train-images-idx3-ubyte.gz

0%| | 0/9912422 [00:00<?, ?it/s]

Extracting mnist_data/MNIST/raw/train-images-idx3-ubyte.gz to
mnist_data/MNIST/raw

Downloading http://yann.lecun.com/exdb/mnist/train-labels-idx1-ubyte.gz
Downloading http://yann.lecun.com/exdb/mnist/train-labels-idx1-ubyte.gz to
mnist_data/MNIST/raw/train-labels-idx1-ubyte.gz

0%| | 0/28881 [00:00<?, ?it/s]

Extracting mnist_data/MNIST/raw/train-labels-idx1-ubyte.gz to
mnist_data/MNIST/raw

Downloading http://yann.lecun.com/exdb/mnist/t10k-images-idx3-ubyte.gz
Downloading http://yann.lecun.com/exdb/mnist/t10k-images-idx3-ubyte.gz to
mnist_data/MNIST/raw/t10k-images-idx3-ubyte.gz

0%| | 0/1648877 [00:00<?, ?it/s]

Extracting mnist_data/MNIST/raw/t10k-images-idx3-ubyte.gz to
mnist_data/MNIST/raw

Downloading http://yann.lecun.com/exdb/mnist/t10k-labels-idx1-ubyte.gz
Downloading http://yann.lecun.com/exdb/mnist/t10k-labels-idx1-ubyte.gz to
mnist_data/MNIST/raw/t10k-labels-idx1-ubyte.gz

0%| | 0/4542 [00:00<?, ?it/s]

Extracting mnist_data/MNIST/raw/t10k-labels-idx1-ubyte.gz to
mnist_data/MNIST/raw

```
/usr/local/lib/python3.7/dist-packages/torchvision/datasets/mnist.py:498:  
UserWarning: The given NumPy array is not writeable, and PyTorch does not  
support non-writeable tensors. This means you can write to the underlying  
(supposedly non-writeable) NumPy array using the tensor. You may want to copy  
the array to protect its data or make it writeable before converting it to a  
tensor. This type of warning will be suppressed for the rest of this program.  
(Triggered internally at /pytorch/torch/csrc/utils/tensor_numpy.cpp:180.)  
    return torch.from_numpy(parsed.astype(m[2], copy=False)).view(*s)
```

4 Define the ANN without convolutional layers for image classification

```
[ ]: class ANNet(nn.Module):
    def __init__(self):
        super(ANNet, self).__init__()
        self.relu = nn.ReLU()
        self.linear1 = nn.Linear(28*28, 500)
        self.linear2 = nn.Linear(500, 100)
        self.linear3 = nn.Linear(100, 10)
    def forward(self, x):
        x = x.view(x.shape[0], -1)
        x = self.relu(self.linear1(x))
        x = self.relu(self.linear2(x))
        x = self.linear3(x)
        return x
```

5 Define a function for validating the model

```
[ ]: def validate(model, data):
    total = 0
    correct = 0
    for i, (images, labels) in enumerate(data):
        images = images.cuda()
        labels = labels.cuda()
        y_pred = model(images)
        value, pred = torch.max(y_pred, 1)
        total += y_pred.size(0)
        correct += torch.sum(pred == labels)
    return correct * 100 / total
```

6 Initialize the neural network and optimizer

```
[ ]: ann = ANNet().cuda()
optimizer = optim.Adam(ann.parameters(), lr=learning_rate)
cross_entropy = nn.CrossEntropyLoss()
```

7 Print the Model Summary

```
[ ]: summary(ann, (1, 224, 224))
```

```
-----
Layer (type)              Output Shape         Param #
-----
```

```

=====
                Linear-1                [-1, 500]                392,500
                  ReLU-2                [-1, 500]                  0
                Linear-3                [-1, 100]                50,100
                  ReLU-4                [-1, 100]                  0
                Linear-5                [-1, 10]                 1,010
=====

Total params: 443,610
Trainable params: 443,610
Non-trainable params: 0
-----

Input size (MB): 0.19
Forward/backward pass size (MB): 0.01
Params size (MB): 1.69
Estimated Total Size (MB): 1.89
-----

```

8 Display the validation accuracy on each epoch

```

[ ]: for n in range(epoch):
      for i, (images, labels) in enumerate(train_data):
          images = images.cuda()
          labels = labels.cuda()
          optimizer.zero_grad()
          prediction = ann(images)
          loss = cross_entropy(prediction, labels)
          loss.backward()
          optimizer.step()
      accuracy = float(validate(ann, val_data))
      print("Epoch:", n+1, "Loss: ", float(loss.data), "Accuracy:", accuracy)

```

```

Epoch: 1 Loss: 0.0039152661338448524 Accuracy: 93.5199966430664
Epoch: 2 Loss: 0.0030660051852464676 Accuracy: 93.5
Epoch: 3 Loss: 0.008405933156609535 Accuracy: 93.75
Epoch: 4 Loss: 0.040073223412036896 Accuracy: 94.37999725341797
Epoch: 5 Loss: 0.015041790902614594 Accuracy: 94.68999481201172
Epoch: 6 Loss: 0.008267774246633053 Accuracy: 96.0999984741211
Epoch: 7 Loss: 0.010446286760270596 Accuracy: 96.30999755859375
Epoch: 8 Loss: 0.000716851616743952 Accuracy: 95.95999908447266
Epoch: 9 Loss: 0.09016045182943344 Accuracy: 95.25999450683594
Epoch: 10 Loss: 0.003279729513451457 Accuracy: 96.72999572753906
Epoch: 11 Loss: 0.003906412981450558 Accuracy: 96.68999481201172
Epoch: 12 Loss: 0.007484755013138056 Accuracy: 95.83999633789062
Epoch: 13 Loss: 0.0004099629877600819 Accuracy: 96.32999420166016
Epoch: 14 Loss: 0.017722729593515396 Accuracy: 96.40999603271484
Epoch: 15 Loss: 0.043866898864507675 Accuracy: 96.56999969482422
Epoch: 16 Loss: 0.0010566012933850288 Accuracy: 96.80999755859375

```

Epoch: 17 Loss: 0.00030694992165081203 Accuracy: 96.72999572753906
Epoch: 18 Loss: 0.01591970957815647 Accuracy: 96.58999633789062
Epoch: 19 Loss: 0.01136217825114727 Accuracy: 96.29000091552734
Epoch: 20 Loss: 0.03559637442231178 Accuracy: 96.29999542236328
Epoch: 21 Loss: 1.262848059013777e-06 Accuracy: 96.62999725341797
Epoch: 22 Loss: 0.0016550994478166103 Accuracy: 96.68000030517578
Epoch: 23 Loss: 0.0001836229785112664 Accuracy: 96.5
Epoch: 24 Loss: 0.0042234198190271854 Accuracy: 96.04000091552734
Epoch: 25 Loss: 0.27225372195243835 Accuracy: 96.72000122070312
Epoch: 26 Loss: 0.008386979810893536 Accuracy: 96.79999542236328
Epoch: 27 Loss: 2.2351693473865453e-07 Accuracy: 96.37999725341797
Epoch: 28 Loss: 0.00573017867282033 Accuracy: 96.5999984741211
Epoch: 29 Loss: 0.0838499516248703 Accuracy: 96.2699966430664
Epoch: 30 Loss: 2.011650508393359e-07 Accuracy: 96.47999572753906