

AI_Assgn_2_Q1

November 10, 2021

1 Import the packages

```
[ ]: import torch
import torchvision
from torch import nn, optim
from torchsummary import summary
```

2 Declare variables for the CNN

- **Epoch** is the number of passes of the entire training dataset through the neural network. A pair of forward and backward propagation indicates a single pass.
- **Batch Size** is the number of samples to work through before updating the weights and biases associated with the model.
- **Learning Rate** controls how much to change the model parameters in response to the prediction error each time the model weights are updated.

```
[ ]: batch_size = 32
epoch = 30
learning_rate = 0.01
```

3 Load the training set and validation set using Dataset and DataLoader

```
[ ]: trans = torchvision.transforms.ToTensor()
train_data = torch.utils.data.DataLoader(
    torchvision.datasets.MNIST(
        'mnist_data', train=True, download=True, transform=trans
    ), batch_size=batch_size
)
val_data = torch.utils.data.DataLoader(
    torchvision.datasets.MNIST(
        'mnist_data', train=False, download=True, transform=trans
    ), batch_size=batch_size)
```

```
Downloading http://yann.lecun.com/exdb/mnist/train-images-idx3-ubyte.gz
Downloading http://yann.lecun.com/exdb/mnist/train-images-idx3-ubyte.gz to
mnist_data/MNIST/raw/train-images-idx3-ubyte.gz
```

```
0%|          | 0/9912422 [00:00<?, ?it/s]
```

```
Extracting mnist_data/MNIST/raw/train-images-idx3-ubyte.gz to
mnist_data/MNIST/raw
```

```
Downloading http://yann.lecun.com/exdb/mnist/train-labels-idx1-ubyte.gz
Downloading http://yann.lecun.com/exdb/mnist/train-labels-idx1-ubyte.gz to
mnist_data/MNIST/raw/train-labels-idx1-ubyte.gz
```

```
0%|          | 0/28881 [00:00<?, ?it/s]
```

```
Extracting mnist_data/MNIST/raw/train-labels-idx1-ubyte.gz to
mnist_data/MNIST/raw
```

```
Downloading http://yann.lecun.com/exdb/mnist/t10k-images-idx3-ubyte.gz
Downloading http://yann.lecun.com/exdb/mnist/t10k-images-idx3-ubyte.gz to
mnist_data/MNIST/raw/t10k-images-idx3-ubyte.gz
```

```
0%|          | 0/1648877 [00:00<?, ?it/s]
```

```
Extracting mnist_data/MNIST/raw/t10k-images-idx3-ubyte.gz to
mnist_data/MNIST/raw
```

```
Downloading http://yann.lecun.com/exdb/mnist/t10k-labels-idx1-ubyte.gz
Downloading http://yann.lecun.com/exdb/mnist/t10k-labels-idx1-ubyte.gz to
mnist_data/MNIST/raw/t10k-labels-idx1-ubyte.gz
```

```
0%|          | 0/4542 [00:00<?, ?it/s]
```

```
Extracting mnist_data/MNIST/raw/t10k-labels-idx1-ubyte.gz to
mnist_data/MNIST/raw
```

```
/usr/local/lib/python3.7/dist-packages/torchvision/datasets/mnist.py:498:
UserWarning: The given NumPy array is not writeable, and PyTorch does not
support non-writeable tensors. This means you can write to the underlying
(supposedly non-writeable) NumPy array using the tensor. You may want to copy
the array to protect its data or make it writeable before converting it to a
tensor. This type of warning will be suppressed for the rest of this program.
(Triggered internally at /pytorch/torch/csrc/utils/tensor_numpy.cpp:180.)
    return torch.from_numpy(parsed.astype(m[2], copy=False)).view(*s)
```

4 Define the CNN for image classification

```
[ ]: class ConvNet(nn.Module):
    def __init__(self):
        super(ConvNet, self).__init__()
        self.conv1 = nn.Conv2d(in_channels=1, out_channels=3, kernel_size=3,
→stride=1)
        self.conv2 = nn.Conv2d(in_channels=3, out_channels=6, kernel_size=3,
→stride=1)
        self.tanh = nn.Tanh()
        self.linear1 = nn.Linear(3456, 10)
    def forward(self, x):
        x = self.tanh(self.conv1(x))
        x = self.tanh(self.conv2(x))
        x = x.view(x.shape[0], -1)
        x = self.linear1(x)
        return x
```

5 Define a function for validating the model

```
[ ]: def validate(model, data):
    total = 0
    correct = 0
    for i, (images, labels) in enumerate(data):
        images = images.cuda()
        labels = labels.cuda()
        y_pred = model(images)
        value, pred = torch.max(y_pred, 1)
        total += y_pred.size(0)
        correct += torch.sum(pred == labels)
    return correct * 100 / total
```

6 Initialize the neural network and optimizer

```
[ ]: convnet = ConvNet().cuda()
optimizer = optim.Adam(convnet.parameters(), lr=learning_rate)
cross_entropy = nn.CrossEntropyLoss()
```

7 Print the Model Summary

```
[ ]: summary(convnet, (1, 224, 224))
```

```
-----
Layer (type)              Output Shape          Param #
=====
```

Conv2d-1	[-1, 3, 222, 222]	30
Tanh-2	[-1, 3, 222, 222]	0
Conv2d-3	[-1, 6, 220, 220]	168
Tanh-4	[-1, 6, 220, 220]	0
Linear-5	[-1, 10]	34,570

Total params: 34,768
Trainable params: 34,768
Non-trainable params: 0

Input size (MB): 0.19
Forward/backward pass size (MB): 6.69
Params size (MB): 0.13
Estimated Total Size (MB): 7.01

8 Display the validation accuracy on each epoch

```
[ ]: for n in range(epoch):
    for i, (images, labels) in enumerate(train_data):
        images = images.cuda()
        labels = labels.cuda()
        optimizer.zero_grad()
        prediction = convnet(images)
        loss = cross_entropy(prediction, labels)
        loss.backward()
        optimizer.step()
    accuracy = float(validate(convnet, val_data))
    print("Epoch:", n+1, "Loss: ", float(loss.data), "Accuracy:", accuracy)
```

```
Epoch: 1 Loss: 0.21218594908714294 Accuracy: 87.68999481201172
Epoch: 2 Loss: 0.15007132291793823 Accuracy: 83.15999603271484
Epoch: 3 Loss: 0.20684774219989777 Accuracy: 86.5
Epoch: 4 Loss: 0.07335351407527924 Accuracy: 87.16999816894531
Epoch: 5 Loss: 0.18528632819652557 Accuracy: 83.43000030517578
Epoch: 6 Loss: 0.22310829162597656 Accuracy: 83.65999603271484
Epoch: 7 Loss: 0.12215633690357208 Accuracy: 87.08999633789062
Epoch: 8 Loss: 0.13716231286525726 Accuracy: 83.90999603271484
Epoch: 9 Loss: 0.13552477955818176 Accuracy: 84.41999816894531
Epoch: 10 Loss: 0.1604619026184082 Accuracy: 84.18999481201172
Epoch: 11 Loss: 0.2963145971298218 Accuracy: 86.19999694824219
Epoch: 12 Loss: 0.12405643612146378 Accuracy: 86.25999450683594
Epoch: 13 Loss: 0.038135409355163574 Accuracy: 84.80999755859375
Epoch: 14 Loss: 0.06628179550170898 Accuracy: 83.58999633789062
Epoch: 15 Loss: 0.44077709317207336 Accuracy: 85.54000091552734
Epoch: 16 Loss: 0.08903875946998596 Accuracy: 84.32999420166016
Epoch: 17 Loss: 0.10942036658525467 Accuracy: 85.45999908447266
```

Epoch: 18 Loss: 0.08876486867666245 Accuracy: 84.68000030517578
Epoch: 19 Loss: 0.0831499993801117 Accuracy: 83.75
Epoch: 20 Loss: 0.029688647016882896 Accuracy: 86.0
Epoch: 21 Loss: 0.0964929386973381 Accuracy: 85.38999938964844
Epoch: 22 Loss: 0.06852761656045914 Accuracy: 85.16999816894531
Epoch: 23 Loss: 0.05163497105240822 Accuracy: 85.15999603271484
Epoch: 24 Loss: 0.04711154103279114 Accuracy: 86.72000122070312
Epoch: 25 Loss: 0.0784594789147377 Accuracy: 86.65999603271484
Epoch: 26 Loss: 0.19556772708892822 Accuracy: 85.07999420166016
Epoch: 27 Loss: 0.034340690821409225 Accuracy: 84.61000061035156
Epoch: 28 Loss: 0.16018325090408325 Accuracy: 84.15999603271484
Epoch: 29 Loss: 0.05336381122469902 Accuracy: 83.63999938964844
Epoch: 30 Loss: 0.12666332721710205 Accuracy: 87.04999542236328